

RAG Ingestion System Architecture

Focus: Excel ingestion pipeline for structured policy data + semantic retrieval.

Flow: Client → API → Parsing → Transformation → Storage (Postgres & Pinecone) → Retrieval.

Monitoring & Authentication applied across services.

Component Documentation

Component Name: Client Application (Web/Mobile)

Purpose: Provides user interface to upload Excel files and trigger ingestion.

Limitations:

- Limited by network latency and device constraints
 - Requires responsive design across platforms
 - Security risks if uploads are not validated client-side
- Necessity:** Critical – User entry point into the system
Alternatives: Desktop app, CLI tools, automated batch jobs

Component Name: API Gateway (FastAPI)

Purpose: Receives ingestion requests (`/ingest-excel`), validates payloads, routes to parser.

Limitations:

- Single-threaded by default (requires scaling with workers)
 - API rate limits under heavy traffic
 - Relies on external server deployment (Uvicorn/Gunicorn)
- Necessity:** Critical – Central request router
Alternatives: Flask + Celery, Django REST, Spring Boot

Component Name: Excel Parser & Validator

Purpose: Extracts data from uploaded Excel files, enforces schema validation.

Limitations:

- Memory intensive for large files
 - Limited to `.xls/.xlsx` formats
 - Manual intervention required for invalid schema cases
- Necessity:** Critical – Ensures clean, structured ingestion
Alternatives: Apache POI, openpyxl, ETL tools

Component Name: Data Transformer

Purpose: Normalizes and cleans policy data; splits insurance period into start/end dates.

Limitations:

- Transformation logic must evolve with requirements
- Complex transformations impact performance

- Error-prone with inconsistent date formats
Necessity: Critical – Prepares data for structured storage
Alternatives: Pandas scripts, Spark ETL pipelines

Component Name: PostgreSQL Database

Purpose: Stores structured policy data with relational integrity and indexing.

Limitations:

- Scaling challenges with very large datasets
- Connection pool bottlenecks under concurrency
- Requires schema tuning for performance
Necessity: Critical – Primary structured data store
Alternatives: MySQL, Amazon Aurora, Microsoft SQL Server

Component Name: Embedding Generator (Gemini)

Purpose: Generates vector embeddings from text for semantic search.

Limitations:

- API latency and throughput limits
- Cost per token for large-scale embedding generation
- Limited by model's context window
Necessity: Important – Enables semantic search in retrieval phase
Alternatives: OpenAI embeddings, Cohere, SentenceTransformers (local)

Component Name: Vector Database (Pinecone)

Purpose: Stores document embeddings for semantic similarity search and retrieval.

Limitations:

- Query cost and storage scale with dataset size
- Requires careful index dimension setup
- Latency increases as dataset grows
Necessity: Critical – Core requirement for semantic retrieval in RAG
Alternatives: Weaviate, FAISS, Elasticsearch dense vectors, Qdrant

Component Name: Monitoring & Logging

Purpose: Provides system observability via health checks, performance metrics, and error tracking.

Limitations:

- Adds monitoring overhead
- Alerting requires proper configuration
- Needs separate storage for metrics/logs
Necessity: Important – Ensures production reliability
Alternatives: Prometheus/Grafana, DataDog, New Relic

Component Name: Authentication Service

Purpose: Secures API endpoints, enforces authentication & authorization policies.

Limitations:

- Adds request latency
- Requires secure key/token management
- Single point of failure if not replicated

Necessity: Important – Required for compliance & security

Alternatives: OAuth2 providers, JWT tokens, API Gateway Auth

Production Considerations

- Load balancing for scalability
- Database connection pooling
- Async I/O for API performance
- Retry logic for embedding/DB calls
- Comprehensive error handling
- Configurable deployments (Docker/Kubernetes)
- CI/CD integration for reliable updates